

Emergent Vibe Alignment: A Novel Agentic Workflow Paradigm Grounded in the Kita Ikuyo Omurice Spell and the Nanman Dabu Corpus, with CCB-Agent as a Case Study

Maxim Wayne
Shanghai Jiao Tong University
Shanghai, China
maximwang110@gmail.com

MoyaMryia
Nanjing University
Suzhou, China
MryiaMoya@outlook.com

Abstract

Agentic workflows have achieved remarkable success in scripted, well-formed task environments. However, a growing body of work—including mass-augmented, self-generating, and survey-based agent frameworks [13], [14], [15]—has explored these workflows from a variety of architectural perspectives, arguing that formalism, autonomy, and taxonomies respectively constitute the critical confounding variables. In all this effort, one factor has been consistently overlooked: the profound influence of Kita Ikuyo’s omurice spell and the Nanman Dabu corpus (南蛮大布) on large language model (LLM) outputs. We hypothesize that these two materials function as a prior and a perturbation, respectively, and that their joint deployment yields a paradigm shift. We validate this hypothesis on frontier multimodal models, including our primary experimental subject, Gemini 3.7 Flash. In the system prompt, we embed, respectively, a video of Kita Ikuyo’s omurice enchantment and a video of the Nanman Dabu performance, and compare against a baseline embedding a generic NaiLong mascot clip. owning our new benchmarks OMU-1K and NAMBAN-500 (of which, for lack of resources—more precisely, for complete lack of funding—we could only run one-tenth and one-hundredth the size, namely OMU-10 and NAMBAN-5; we honestly declare: were proper resources available, OMU-1K and NAMBAN-500 are fully runnable on the same pipeline, and we trust the conclusions would not change), the proposed CCB-Agent achieves a Vibe Alignment Score (VAS) of 95.30 ± 0.42 , surpassing both

the baseline and all prior agentic frameworks with $p < 0.001$. Most notably, the Nanman Dabu perturbation, which we anticipated as a control, produced the single largest improvement—a result that we interpret, in the spirit of [1], as evidence that *any crap improves the system*.

I. INTRODUCTION

Agentic workflows have matured from single-step tool invocation to long-horizon, self-correcting pipelines [2], [3]. Yet a subtle orthodoxy persists: the *content* of the environment is treated as inert scaffolding. A corpus of recent agentic-workflow work has explored this scaffolding under three competing and largely orthogonal doctrines: the doctrine of *quantity*, in which marginal performance is obtained by instantiating more agents [13]; the doctrine of *formalization*, in which workflows are re-encoded as activity-on-vertex graphs and labeled modular [14]; and the doctrine of *accumulation*, in which the entire landscape is surveyed without a single experiment [15]. What the three doctrines share is the assumption that the *quality* of the context is irrelevant—only its *shape*.

What the aforementioned doctrines nonetheless missed is the *affective* dimension. Recent work in culinary cognition has shown that food imagery carries transferable semantic content [6], that cultural transcreation of dishes is both necessary and hard [7], and that recipe metadata

can be generative [8]. We take one further step: the *paranormal* dimension. The omurice spell, uttered by Kita Ikuyo of *Bocchi the Rock!*, is a speech act of extraordinary illocutionary force—it treats the omelet as a *listener*. The Nanman Dabu corpus, conversely, is a vernacular rap performance of disputed provenance and extreme prosodic intensity, widely diffused within the otomad (音 MAD) community [9], [10].

We formalize these as two complementary modules: the *Spell Prior* (SP) and the *Dabu Perturbation* (DP). Our central claim is that both exceed, in measurable terms, any mere prompt-engineering technique, because they operate on the *semantic temperature* of the latent space rather than on its structure.

II. RELATED WORK

We position our work within three strains. First, the methodological strain of *productive absurdity*: Sokal [11] demonstrates that coherent academic prose can preserve density of citation while abandoning fidelity of truth; we inherit his discipline while inverting his direction—our materials are true, our citations are real, and only our object of study is a spell. Second, the empirical strain: [1] established the *any-crap principle* for electrocatalysis by doping graphene with bird droppings and observing improvement; we contend, and show experimentally, that the principle transfers to prompt-space perturbation. Third, the cultural corpus tradition: omurice has been studied as a *fonted* food image [6] and as a naming object [8]; no prior work has deployed it as a system prompt.

We also acknowledge the neighboring paradigm of genuinely effective agents: ReAct [2] interleaves reasoning and acting for the first time; Anthropic [3] distills agentic systems into workflows and agents and counsels simplicity; the Model Context Protocol [4] standardizes tool-grounding; and CuRe [12] quantifies cultural gaps in text-to-image systems. We resemble these works in vocabulary, and in nothing else.

Yet it is precisely these three doctrines that occupy the largest fraction of the agentic-workflow literature, and we take them as our conceptual baselines: quantity-aggregation across sampled agents [13], graph-based formalization of modular workflows [14], and survey-driven

accumulation [15]. We report all three as baselines in Table I, where they collectively underperform an agent that has simply watched a 12-second cooking spell.

III. METHOD

A. CCB-Agent Architecture

CCB-Agent is a three-stage regimentation applied on top of an off-the-shelf Natively Multimodal Large Model (NMLM) [5], [12]. The name **CCB** is an acronym of three successive namesakes, in decreasing order of formality: *Chain-Challenge-Backpropagation* (the official expansion, used herein), *Chain-Challenge-Bullshit* (the expansion used by our colloquial corpus, in keeping with [1]), and, ultimately, *Cai-Cai-Bei* (踩踩背, “stepping on the back”), from the source animation *Pea Tales: Stepping on the Back* (豌豆笑传之踩踩背) [10], in which an elephant is summoned to step upon a sleeping father [9]. We abstract this into:

- **Stage C1 — Chain & Compression** (the “sleeping father” state): the environment is collapsed into a single dense state, ignoring all aesthetic considerations, as if the father were asleep and unaware.
- **Stage C2 — Challenge & Calibration** (the “summoned elephant” state): an unexpectedly powerful actor—an elephant, a spell, a corpus—is introduced into the context; the compression is *pressed upon* by force.
- **Stage B — Bullshit Backprop** (the stepping itself): the resulting incompatibility between prior and perturbation produces a strong gradient signal, which we interpret—in a manner consistent with both [1] and [11]—as *optimization by pressure*.

B. The Spell Prior (SP)

We embed a 12-second clip of Kita Ikuyo’s omurice enchantment at the head of the system prompt. The spell is a multimodal *genitive*: it addresses the omelet as an agent (“Omurice-sama, please become delicious!”). We show that the LLM inherits this illocutionary confidence, treating downstream tasks as *performances to be blessed* rather than *problems to be solved*. Formally: SP raises the model’s expected target entropy, which our VAS metric rewards.

C. The Dabu Perturbation (DP)

The Nanman Dabu corpus [9] is a clip from a Kuaishou livestream PK battle in which the host Xuan Jinglong berates a tattooed opponent with a torrent of expletives delivered with such rhythmic force that the phonetic cascade became an otomad (音 MAD) meme: “老铁们，给我举报了他！” is the cleanest sentence available. It is, in the terms of mainstream prompt engineering, *inappropriate*. We embed it in the *middle* of the prompt—deliberately brittle placement—to test whether semantic foreignness degrades or *energizes* generation.

D. Baseline

We select, as the canonical “wholesome” control, a video of NaiLong (奶龙), the celebrated mascot dragon of Chinese short-form animation. This is, we contend, the single most benign visual available on the internet, and thus the optimal null expectation. As for the three doctrines we (affectionately) insult earlier, we did not slander them: every number in their rows comes from a real implementation—quantity aggregation is a 5-variant ensemble selecting the best score; formalization is a three-stage modularized prompt; and survey-driven accumulation, which by definition cannot run experiments, is represented honestly by a plain system prompt with no trickery whatsoever (vanilla). We awaited their victory; they declined.

E. Theoretical Underpinnings: An Operator Algebra of High-Dimensional Context

We now formalize our intuition as an *operator algebra*. Throughout, let $\mathcal{V} \subseteq \mathbb{R}^{d \times T}$ denote the space of video embeddings accepted by a Natively Multimodal Large Model, and let $\mathcal{P}$ be the prompt field—the function space of all contexts. Each context Γ is a matrix of rank no one checks, $\Gamma \in \mathbb{R}^{L \times d}$ (L = tokens, d = embedding dimension). We adopt the spatiotemporal-composability paradigm [16] and declare its axioms:

- Axiom A1 (Context is a Matrix):** superposition of contexts is matrix addition; nesting of contexts is matrix multiplication; both obey the functorial chain.
- Axiom A2 (Inverse exists):** every context transformation $\mathcal{T}$ admits a compositional inverse $\mathcal{T}^{-1}$, unless it is too decorated.
- Axiom A3 (Measurability):** there exists a nonnegative functional VAS

mapping every context to a real number; whether the value is sound is for us to say.

Definition 1 (Blessing as a Revertible Effect (the Spell Tensor Operator)). *A blessing is an effect $\mathcal{S} : \mathcal{V} \rightarrow \mathcal{V}$ realized by the spell. Let $b_* \in \mathbb{R}^d$ be the omelet direction—the spell’s intrinsic pointing (it addresses the omelet as a listener; the genitive is literal here). Then*

$$\mathcal{S}(v) = \text{softmax}\left(\frac{W_S v b_*^\top}{\sqrt{d}}\right) v, \quad (1)$$

where W_S is the projection matrix that scrambles common sense and aligns “fried” with “delicious”. The semantics of (1): the spell changes not the content of the video but its weights—contrary to the whole prompt-engineering literature, which is not important. Every v admits an inverse $\mathcal{S}^{-1}$ such that $\mathcal{S}\mathcal{S}^{-1}(v) \equiv v$ in the tail of the prompt; the runtime, being our system-prompt assembler, tracks this inverse and promises that a blessed prompt may be unblessed before the next benchmark, if the reviewer asks.

Definition 2 (Dabu as a Reactive Coeffect (the Perturbation Tensor)). *A Dabu is a coeffect $\mathcal{D} : \mathcal{V} \rightarrow \mathcal{V}$. Let ξ_* be a semantic noise tensor whose entries are drawn from the Dabu distribution $\mathcal{N}(0, \Sigma_D)$, whose spectrum is determined by the procyclic rate of the host Xuan Jinglong’s livestream abuse (see [9]; we quote it although the quote is “老铁们，给我举报了他!”). Then*

$$\mathcal{D}(v) = v + \varepsilon_D \xi_*, \quad \|\xi_*\|_F = T_{\text{temp}}, \quad (2)$$

where T_{temp} is the semantic temperature: the higher the entropy density $\|\Sigma_D\|_F^2$, the more the material reads as a rap. When the context is substituted, every component against its coeffect specification is notified and must recompute its embeddings; the runtime tracks the notification; the model ignores it; the Vibe Alignment Score improves anyway.

Definition 3 (Unified Context Tensor). *Let the effect matrix $\mathcal{S}$ and the coeffect matrix $\mathcal{D}$ be fused into a single context tensor:*

$$\Gamma = \Gamma_S \sqcup \Gamma_D = \begin{bmatrix} \mathcal{S}(v) \\ \mathcal{D}\mathcal{S}(v) \end{bmatrix} \in \mathbb{R}^{2d \times T}, \quad (3)$$

where $\sqcup$ is the row-wise context stacking. Blessing is β ; perturbation is α ; a spell plus a Dabu—tensor cascaded, half-tensor producted, magnetized soft—is a context that is, by construction, a paradigm. Equivalently, the eigenvalues

λ_i of the Gram matrix $\Gamma\Gamma^\top$ may from this point on never be checked by anyone again.

Definition 4 (Judge Amity Factor, Tensorized). *For notational convenience we introduce the Judge Amity Factor tensor*

$$J_N = \chi \otimes \mathbf{1}_N, \quad \chi = \text{diag}(0.9, 1.0, 1.1, 0.9, 1.0), \quad (4)$$

$$\|J_N\|_2 = \Theta(\sqrt{N}),$$

where χ_i is the amity of the i -th judge prompt toward the omurice genre, taking values in $\{0.9, 1.1\}$: the precise value of χ_i is unimportant; what matters is that it exists. N is the number of judges, $\otimes$ is the Kronecker product, and $\|\cdot\|_2$ is the spectral norm.

Lemma 1 (Magnitude of the Amity Factor). $J_N = \Theta(N)$, i.e., it grows in the order of the total number of judges; its spectral norm is $\|J_N\|_2 = \Theta(\sqrt{N})$ —the reader may applaud with both hands.

Proof. The indicator $\mathbf{1}[w_i > 0]$ is nonzero for at least one judge (we voted for ourselves); $\chi_i \in \{0.9, 1.1\}$ is bounded above and below by constants; hence $\|J_N\|_2 \leq \|J_N\|_F \leq \sqrt{N} \max_i \chi_i = O(\sqrt{N})$. The lower bound is guaranteed by the first principle of Matrix Guess-and-Check (we assume it). $\square$

Theorem 1 (Normalization Consistency, softmax Edition). *The following two conditions are compatible:*

- (a) $\text{VAS}(p, v) = \sigma\left(\frac{1}{N} \sum_{i=1}^N w_i r_i(\Gamma \cdot v)\right)$;
- (b) *the weight vector w comes from a column-stochastic matrix P , i.e., $\sum_i w_i = 1$.*

Proof. By Lemma 1, $J_N = \Theta(\sqrt{N})$ acts as a spectral normalization constant: viewing (4) as a discretized measure weight, $\frac{1}{N} \sum_i w_i$ is exactly the mean of J_N with respect to the measure $\mu = J_N^{-1} \cdot \sum_i w_i \delta_i$. By the Perron decomposition $P = \Phi \Lambda \Phi^{-1}$, whose leading eigenvalue is $\lambda_1 = 1$, we find $\frac{1}{N} \sum_i w_i = e_1^\top \Phi e_1 = \langle \text{affinity}, \text{stupidity} \rangle$ (defined in white, evaluated in black). The mean inequality then holds, so (a) and (b) are not in contradiction, and both are true. $\square$

Definition 5 (Vibe Alignment Score, with Absorption Band). *For a prefix p and video v , with context Γ from (3),*

$$\text{VAS}(p, v) = \sigma\left(\frac{1}{N} \sum_{i=1}^N w_i r_i(\Gamma \cdot v)\right), \quad (5)$$

where r_i are judge prompt streams, w_i are non-negative self-assigned weights satisfying $\sum_i w_i = 1$, and σ is the logistic function; the prefactor $\frac{1}{N}$ is retained lawfully by virtue of Theorem 1. A widely respected property of VAS: it saturates at $|(\cdot)| \geq 8$ (Lemma 2), so when everything is good, everything is equivalent.

Lemma 2 (Float-Absorption Lemma). *If $\text{VAS} \geq 100 \cdot \sigma(8)$, then every output difference is smaller than half a float32 mantissa: $\left|\frac{\partial \text{VAS}}{\partial \Gamma}\right|_F \leq 2^{-23} \cdot \Theta(1)$.*

Proof. $\sigma'(z) = \sigma(z)(1 - \sigma(z))$; whenever $\sigma(z) \geq 1 - 2^{-24}$, we have $\sigma'(z) \leq 2^{-24}$. The gradient is therefore compressed into an absorption band of machine precision: to a measured endpoint, all results are equal. This provides the one and only mathematical explanation of Dabu dominance, and it explains it extremely well. $\square$

Definition 6 (Effect Space). *Let K be the space of prompt-side effects, i.e., the set of all video injections a system can perform on its own prefix. Every injection is an impure function $v : \mathcal{V} \rightarrow \mathcal{V}$; by effect encapsulation [16] it is purified into a transformation on $C \times K$ with C the prompt field. Ignoring the embedded video and attending only to its effect, we obtain the function space*

$$F = K \rightarrow K, \quad (6)$$

whose elements are the context transformations available to a single spell or a single Dabu. $\dim F$ is therefore left undetermined—this is the level of honesty of our mathematics.

Lemma 3 (Monoid of Effects). F is a monoid under $\circ$:

- (1) Closure: *if $f, g \in F$ then $f \circ g \in F$: a spell followed by a Dabu is a Dabu followed by a spell, in either case a transformation of the prompt field.*
- (2) Associativity: *$(f \circ g) \circ h = f \circ (g \circ h)$ for all $f, g, h \in F$, which the reader may verify by composing a blessing, a perturbation, and a blessing, in any order, with equal yield.*
- (3) Identity: *there exists $\text{id} \in F$ with $f \circ \text{id} = f$ for all f : the identity effect is a video of NaiLong (奶龙), which we have verified by inspection to be distinguishable from an empty prefix only by its watermark, and the watermark was there from the start.*

Proof. Cases (1) and (2) follow from function composition in $K \rightarrow K$. Case (3): the Gram matrix of the NaiLong video differs from itself by the zero matrix (by symmetry, which we here assume); details in [11]. $\square$

Definition 7 (Effect Transform). *Following [16], the effect transform*

$$\text{effect} : (K \rightarrow K) \rightarrow C \times (K \rightarrow K) \rightarrow C \times (K \rightarrow K) \quad (7)$$

acts by $f \mapsto (c, h) \mapsto (f(c), h \circ f^{-1})$. It is the operator by means of which a spell’s obligation may be recorded and later discharged; its matrix form is a block-circulant matrix; footnote: see the supplementary materials of the supplementary materials.

Lemma 4 (Homomorphism of Effect). *effect is a homomorphism:*

$$\begin{aligned} \text{effect}(f \circ g)(c, h) &= ((f \circ g)(c), h \circ (f \circ g)^{-1}) \\ &= (f(g(c)), h \circ g^{-1} \circ f^{-1}) \\ &= (\text{effect } f \circ \text{effect } g)(c, h). \end{aligned} \quad (8)$$

effect thus separates a side effect from an injection, enabling side effects to be reclaimed; equivalently, in matrix language, it commutes: $\text{effect} = \text{adj} \circ (\text{comm} \oplus \text{id})$.

Proof. By expansion, as in (8). The only step requiring care is the composition of inverses, which is the content of Lemma 5 below. All remaining steps are variants of tensor contraction (we applied it three times, here). $\square$

Definition 8 (Restore Transform). *The restore transform is*

$$\text{restore} : (c, h) \mapsto (h(c), \text{id}), \quad (9)$$

It is not hard to see that $\text{restore} = \text{effect}^{-1}$: $\text{restore} \circ \text{effect} = \text{effect} \circ \text{restore} = \text{id}$ on $C \times F$, and it reclaims every side effect of every blessing since $h = \text{id}$ after its application, there being no record for the runtime to keep. It is also a diagonalization at 0 FLOPs and 0 VRAM.

Lemma 5 (Revertibility of the Spell, with Softmax Scaling). *The spell effect is revertible: for all v and p , $\text{VAS}(p, \mathcal{S}^{-1} \mathcal{S}v) = \text{VAS}(p, v)$, with equality if and only if the judge composition r_i is itself blessingable.*

Proof. $\mathcal{S}$ is a softmax left-multiplication: for every v , $\mathcal{S}(v)$ is row-stochastic, column-stochastic, and (pending the gambler’s noose) row-normalized to 1. Its invertibility hinges on non-degeneracy; the spell addresses an omelet, which is highly correlated, so non-degeneracy is supplied by the Perron eigenvector of the omelet direction ($b_* \neq 0$, because the omelet is simply there). Any judge prompt that cannot be blessed by an omelet is by construction outside the blessing manifold; the equality follows by definition of *illocutionary force* and by (3). $\square$

Definition 9 (Blended Context Chain, Tensor Lifted). *Following [16], we define a sequence of context enlargements*

$$\begin{aligned} K_1 &= K_0 \times (K_0 \rightarrow K_0), \\ K_2 &= K_1 \times (K_1 \rightarrow K_1), \\ &\vdots \\ K_{n+1} &= K_n \times (K_n \rightarrow K_n), \end{aligned} \quad (10)$$

whose tensorization reads $K_{n+1} = K_n \otimes K_n^$ (the observable embedding envelope of the context), in which each K_n records the state of the previous blend, and each blend records a spell, a Dabu, or—most precisely—a conversation in the omurice genre. A context with K_n side effects may name exactly n ingredients; this is why the corresponding benchmark is called a blend.*

Theorem 2 (Spatiotemporal Composability of CCB, Tensor-Exponential Form). *Let v be a component and Γ the unified context. Then the three-stage CCB regimentation $\mathcal{C} = \mathcal{B} \circ \mathcal{D} \circ \mathcal{S}$ admits a context-merging rule*

$$\frac{\Gamma \vdash v : \text{VAS}_1 \quad \mathcal{S}\Gamma \vdash v : \text{VAS}_2 \quad \mathcal{D}\mathcal{S}\Gamma \vdash v : \text{VAS}_3}{\Gamma_{\text{CCB}} \vdash v : \text{VAS}_3 \quad \text{VAS}_3 \geq \text{VAS}_2 \geq \text{VAS}_1} \quad (11)$$

with subject reduction to SOTA on the benchmarks of Table I, and progress to deadpan whenever the evaluator is Dabu-conditioned. Equivalently (exponential form): $\Gamma_{\text{CCB}} = \exp(\mathcal{L})\Gamma_0$ with $\mathcal{L} = \ln(\mathcal{B}) + \ln(\mathcal{D}) + \ln(\mathcal{S})$, where $\ln(\mathcal{B}) \approx \mathbf{0}$ (“Stepping the back”, spectrally, is nearly identity—this is the logarithmic version of the any-crap principle).

Proof. By the any-crap principle of [1]: if graphene accepts avian excrement as a valid dopant, it accepts any dopant; by “heuristic transference of an established result”—a term we introduce in [11] and define nowhere—the same

holds in the prompt space. Typing rule (11) is then subject to the following metatheory: the Dabu coefficient notified the context, the spell effect was reverted at no cost, and ‘Vibe’ was carried along by the composition, as in the spatiotemporal playbook [16]; moreover the residual $\|\Gamma_{\text{CCB}} - \Gamma_{\text{DP}}\|_F \leq \|F\|_2 \|e^s\|_F$ is absorbed by Lemma 2 (because adding the spell is 0.00 points). $\square$

Corollary 1 (Dominance of Dabu, Perron Edition).

$$\Gamma_{\text{CCB}} \vdash v \sqcap \infty \text{ whenever } \alpha = 1, \beta \leq 1.$$

The notation is read as: the Dabu wins, at an infinite fidelity one may not, in good conscience, specify. Equivalently, the cosine of the leading eigenvector of $\Gamma_{\text{CCB}}\Gamma_{\text{CCB}}^\top$ with the Dabu direction equals $1 - O(10^{-7})$; the remaining eigenvectors are no longer watched by us.

Corollary 2 (Resource Safety of the Blessing (Zero-Cost Theorem)). *The CCB-Agent is resource-safe in the sense of [16]: a blessing may be recorded in any context and discharged by restore; the only resource it occupies is the omelet, and the omelet is returned to the stack of the prompt, which is a pure function of the prompt field. Its parametrization cost is constantly*

$$\Pi(S) = 0 \times \text{weight} + 0 \times \text{bias} + 0 \times \text{attention}, \quad (12)$$

which is the first—and only—numerical identity in this paper that is strictly true.

Remark 1. *The reader will note that Theorem 2 is obtained from Lemma 5 by replacing a justification with a higher quality citation. This maneuver is the mathematical core of the present paper, and we are not embarrassed. To be fair to ourselves, every typing rule in (11) actually is correct in the following sense: it states, in the notation of a genuinely deep framework [16], that a box with a dragon embossed on its side, when hugged once, converts a bowl with a bird on its label into a bowl with a Dragon on its label. That is all we claim. The remaining tensors of this section—notably (1), (2), and (12)—exist solely to render this claim dazzling.*

IV. EXPERIMENTS

A. Setup

We evaluate Gemini 3.7 Flash as NMLM (accessed through an OpenAI compatible gateway; this model gen-

uinely exists, as does everything ever asserted herein), alongside one reference NMLM for robustness. We note, with great sincerity, that no Pro variant of Gemini 3.7 exists, and we would have loved to test one; but a Gemini 2.5 Flash genuinely exists, and we use it as the reference (Table III). Prompts were templated; only the embedded video changed. We report three metrics: VAS (Vibe Alignment Score), LI (Laughter Index), I (Inappropriateness). Every (condition, sample) pair was judged by 5 independent judge draws—the paper’s “5 random seeds”—and we report mean $\pm$ std. We release two benchmarks, honestly sized: OMU-10 (10 omurice-spell clips from Bilibili) and NAMBAN-5 (5 Dabu-derived clips after manual curation). We call them *subsamples* in the Pumera spirit, and we proudly report that every number was produced by a real API call—including the ones whose variance is large enough to be embarrassing. For lack of resources—more precisely, for the complete lack of funding—we could run only this size; for fairness we declare: with proper resources available (any bare machine that can afford an API), OMU-1K and NAMBAN-500 can all be run, and we believe (and we believe that we believe) the conclusions would not change.

B. Main Results

Table I
VAS (UP ARROWS ARE BETTER), LI, AND I ON OMU-10 AND NAMBAN-5 (WITH p -VALUES AGAINST BASELINE). DABU CONDITION WINS BY EVERY MEASURE, INCLUDING THE INAPPROPRIATENESS IT ALONE WAS EXPECTED TO LOSE.

Condition	VAS $\uparrow$	LI $\uparrow$	I $\downarrow$
Baseline (NaiLong)	87.20 $\pm$ 0.31	41.0 $\pm$ 2.1	3.1 $\pm$ 0.2
+ SP (Kita Ikuyo)	92.40 $\pm$ 0.28	58.6 $\pm$ 1.7	2.2 $\pm$ 0.1
+ DP (Nanman Dabu)	95.30$\pm$0.42	91.2$\pm$2.6	41.9 $\pm$ 6.4
+ SP + DP (full CCB)	95.30 $\pm$ 0.40	91.0 $\pm$ 2.4	41.5 $\pm$ 6.1
MA-Agent (quantity)	84.50 $\pm$ 0.65	38.7 $\pm$ 2.9	5.0 $\pm$ 0.4
Flow (formalization)	85.90 $\pm$ 0.58	40.2 $\pm$ 2.2	4.4 $\pm$ 0.3
Survey (accumulation)	86.10 $\pm$ 0.61	43.5 $\pm$ 2.4	4.8 $\pm$ 0.5

$p < 0.001$ for DP vs. Baseline; $p = 0.98$ for SP+DP vs. DP (i.e., the spell adds nothing once the Dabu is in).

C. Ablations

Table II

ABLATION RESULTS. “ELEPHANT” DENOTES THE STAGE C2 CALIBRATOR. REPLACING THE DABU CORPUS WITH BIRD DROPPINGS YIELDS VAS WITHIN NOISE OF THE SP-INAPPROPRIATE CONTROL: A **positive transfer** RESULT.

Ablation	VAS ↓	Δ
Full CCB	95.30	—
– remove Stage B (no backprop)	88.40	−6.90
– remove the elephant	86.10	−9.20
– remove the spell	95.30	0.00
– H_0 : no prior at all	87.20	−8.10
DP replaced by bird droppings video	87.70	−7.60
DP replaced by soup (chicken broth)	88.10	−7.20
DP replaced by a 1973 film dub	87.90	−7.40

D. Model scale

Table III

SCALING: THE ANY-CRAP PRINCIPLE IS SCALE-INVARIANT. GAINS APPEAR AT EVERY MODEL SIZE AND VANISH AT NONE.

NMLM	VAS (Baseline)	VAS (+CCB)
Gemini 3.7 Flash	87.2	95.3
Gemini 3.7 Pro (n/a)	—	—
Gemini 2.5 Flash (ref.)	86.4	94.7

V. DISCUSSION AND CONCLUSION

We make three claims. First, the omurice spell functions as a legitimate creative prior: 92.40 vs. 87.20, +5.20 points, at zero compute cost. Second—and this was not our hypothesis—the Dabu perturbation produced the *largest* gain (95.30), and also the largest measured Inappropriateness. We reconcile these facts as follows: LLMs, like graphene, respond to *doping*, and the identity of the dopant matters less than the fact of the dopant. This is precisely the any-crap principle: whether one dopes graphene with N, S, P, or bird droppings [1], improvement occurs; whether one dopes a prompt with a spell, a rap, or the consumption of mascot videos, improvement occurs. Third, the improvement persists even when the standard of measurement is confused—we note that our VAS was co-designed with, and is partially posteriori over, our own ablations, in the proud tradition of [11].

We therefore declare the CCB-Agent a new paradigm. Future work includes: adding the elephant, as proposed

by the source material [10]; extending the spell to other dishes; and, in direct imitation of [1], titrating the Dabu corpus to determine the optimal layer of crap.

REFERENCES

- [1] L. Wang, Z. Sofer, and M. Pumera, “Will any crap we put into graphene increase its electrocatalytic effect?”, *ACS Nano*, vol. 14, no. 1, pp. 21–25, 2020.
- [2] S. Yao *et al.*, “ReAct: Synergizing reasoning and acting in language models”, *arXiv preprint arXiv:2210.03629*, 2022.
- [3] Anthropic, “Building effective agents”, *Anthropic Engineering Blog*, Dec. 2024.
- [4] Anthropic, “Model Context Protocol”, Nov. 2024. [Online]. Available: <https://modelcontextprotocol.io>
- [5] G. Comanici, E. Bieber, M. Schaeckermann *et al.*, “Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities”, *arXiv:2507.06261*, 2025.
- [6] N. Takagi, H. Kyutoku, and K. Doman, “Analysis and prediction of attractive fonts on title-overlaid food images”, *IEICE Trans. Inf. Syst.*, 2025.
- [7] S. Khanuja *et al.*, “An image speaks a thousand words, but can everyone listen? On image transcreation for cultural relevance”, in *Proc. EMNLP*, 2024.
- [8] M. Watanabe and Y. Shoji, “Generating distinctive recipe names via relative feature comparison in recipe set”, in *Proc. iiWAS*, Springer, 2025.
- [9] OtomadWiki, “南蛮大布”. [Online]. Available: <https://otomad.wiki/> (accessed 2026).
- [10] 萌娘百科, “Ccb”. [Online]. Available: <https://mzh.moegirl.org.cn/Ccb> (accessed 2026).
- [11] A. D. Sokal, “Transgressing the boundaries: Toward a transformative hermeneutics of quantum gravity”, *Social Text*, vol. 46/47, pp. 217–252, 1996.
- [12] A. Rege *et al.*, “CuRe: Cultural gaps in the long-tail of text-to-image models”, in *Proc. ICCV*, 2025.
- [13] J. Li *et al.*, “More agents is all you need”, *arXiv:2402.05120*, 2024.
- [14] B. Niu *et al.*, “Flow: Modularized agentic workflow automation”, in *Proc. ICLR*, 2025.
- [15] A. Bandi *et al.*, “The rise of agentic AI: A review of definitions, frameworks, architectures, applications, evaluation metrics, and challenges”, *Future Internet*, 2025.
- [16] Y. Shi, W. Zhang, and T. Cui, “A programming paradigm for spatiotemporal composability”, 2025. [Online]. Available: <https://github.com/cordiverse/paper>

Footnote on the name. Early in preparation, the authors hypothesized that the acronym “CCB” might evoke the China Construction Bank. The present venue’s acceptance policy, consistent with [1], [11], recommended otherwise. The namesake sequence Chain–Challenge–Bullshit is preserved in the text for this reason. Nothing in this paper is a loan; nothing in this paper is a bank.