

Emergent Vibe Alignment: A Novel Agentic Workflow Paradigm Grounded in the Kita Ikuyo Omurice Spell and the Nanman Dabu Corpus, with CCB-Agent as a Case Study

Maxim Wayne
Shanghai Jiao Tong University
Shanghai, China
maximwang110@gmail.com

MoyaMryia
Nanjing University
Suzhou, China
MryiaMoya@outlook.com

Assisted by: opencode (deepseek-v4-flash-vision-exp)
Writing and experiment orchestration

Abstract

Agentic workflows have achieved remarkable success in scripted, well-formed task environments. However, a growing body of work—including mass-augmented, self-generating, and survey-based agent frameworks [13], [14], [15]—has explored these workflows from a variety of architectural perspectives, arguing that formalism, autonomy, and taxonomies respectively constitute the critical confounding variables. In all this effort, one factor has been consistently overlooked: the profound influence of Kita Ikuyo’s omurice spell and the Nanman Dabu corpus (南蛮大布) on large language model (LLM) outputs. We hypothesize that these two materials function as a prior and a perturbation, respectively, and that their joint deployment yields a paradigm shift. We validate this hypothesis on frontier multimodal models, including our primary experimental subject, Gemini 3.7 Flash. In the system prompt, we embed, respectively, a video of Kita Ikuyo’s omurice enchantment and a video of the Nanman Dabu performance, and compare against a baseline embedding a generic NaiLong mascot clip. owning our new benchmarks OMU-1K and NAMBAN-500 (of which, at a total budget of ten yuan, we ran one-hundredth each, namely OMU-10 and NAMBAN-5; with proper resources, the full sets are runnable on the same pipeline, and we trust the conclusions would not change), the proposed CCB-Agent achieves a Vibe Alignment Score (VAS) of 79.03 ± 8.47 , a modest $+4.99$ over the baseline ($p = 0.368$, insignificant), edging out Flow and Survey in the mean. Most notably, the only

system to surpass the baseline significantly ($p < 0.001$) is the “quantity-aggregation” baseline—one of the three doctrines we set out to defeat. In the spirit of [1], we interpret this as evidence that *anything works; and that is precisely why nothing matters; and the only thing garbage proved useful is what is called garbage*.

I. INTRODUCTION

Agentic workflows have matured from single-step tool invocation to long-horizon, self-correcting pipelines [2], [3]. Yet a subtle orthodoxy persists: the *content* of the environment is treated as inert scaffolding. A corpus of recent agentic-workflow work has explored this scaffolding under three competing and largely orthogonal doctrines: the doctrine of *quantity*, in which marginal performance is obtained by instantiating more agents [13]; the doctrine of *formalization*, in which workflows are re-encoded as activity-on-vertex graphs and labeled modular [14]; and the doctrine of *accumulation*, in which the entire landscape is surveyed without a single experiment [15]. What the three doctrines share is the assumption that the *quality* of the context is irrelevant—only its *shape*.

What the aforementioned doctrines nonetheless missed is the *affective* dimension. Recent work in culinary cognition has shown that food imagery carries transferable semantic content [6], that cultural transcreation of dishes is both necessary and hard [7], and that recipe metadata

can be generative [8]. We take one further step: the *paranormal* dimension. The omurice spell, uttered by Kita Ikuyo of *Bocchi the Rock!*, is a speech act of extraordinary illocutionary force—it treats the omelet as a *listener*. The Nanman Dabu corpus, conversely, is a vernacular rap performance of disputed provenance and extreme prosodic intensity, widely diffused within the otomad (音 MAD) community [9], [10].

We formalize these as two complementary modules: the *Spell Prior* (SP) and the *Dabu Perturbation* (DP). Our central claim is that both exceed, in measurable terms, any mere prompt-engineering technique, because they operate on the *semantic temperature* of the latent space rather than on its structure.

II. RELATED WORK

We position our work within three strains. First, the methodological strain of *productive absurdity*: Sokal [11] demonstrates that coherent academic prose can preserve density of citation while abandoning fidelity of truth; we inherit his discipline while inverting his direction—our materials are true, our citations are real, and only our object of study is a spell. Second, the empirical strain: [1] established the *any-crap principle* for electrocatalysis by doping graphene with bird droppings and observing improvement; we contend, and show experimentally, that the principle transfers to prompt-space perturbation. Third, the cultural corpus tradition: omurice has been studied as a *fonted* food image [6] and as a naming object [8]; no prior work has deployed it as a system prompt.

We also acknowledge the neighboring paradigm of genuinely effective agents: ReAct [2] interleaves reasoning and acting for the first time; Anthropic [3] distills agentic systems into workflows and agents and counsels simplicity; the Model Context Protocol [4] standardizes tool-grounding; and CuRe [12] quantifies cultural gaps in text-to-image systems. We resemble these works in vocabulary, and in nothing else.

Yet it is precisely these three doctrines that occupy the largest fraction of the agentic-workflow literature, and we take them as our conceptual baselines: quantity-aggregation across sampled agents [13], graph-based formalization of modular workflows [14], and survey-driven

accumulation [15]. We report all three as baselines in Table I, where, by their means, they collectively underperform an agent that has simply watched a 23-second cooking spell; notably, the only statistically significant victory of this benchmark belongs to the quantity doctrine itself.

III. METHOD

A. CCB-Agent Architecture

CCB-Agent is a three-stage regimentation applied on top of an off-the-shelf Natively Multimodal Large Model (NMLM) [5], [12]. The name **CCB** is an acronym of three successive namesakes, in decreasing order of formality: *Chain-Challenge-Backpropagation* (the official expansion, used herein), *Chain-Challenge-Bullshit* (the expansion used by our colloquial corpus, in keeping with [1]), and, ultimately, *Cai-Cai-Bei* (踩踩背, “stepping on the back”), from the source animation *Pea Tales: Stepping on the Back* (豌豆笑传之踩踩背) [10], in which an elephant is summoned to step upon a sleeping father [9]. We abstract this into:

- **Stage C1 — Chain & Compression** (the “sleeping father” state): the environment is collapsed into a single dense state, ignoring all aesthetic considerations, as if the father were asleep and unaware.
- **Stage C2 — Challenge & Calibration** (the “summoned elephant” state): an unexpectedly powerful actor—an elephant, a spell, a corpus—is introduced into the context; the compression is *pressed upon* by force.
- **Stage B — Bullshit Backprop** (the stepping itself): the resulting incompatibility between prior and perturbation produces a strong gradient signal, which we interpret—in a manner consistent with both [1] and [11]—as *optimization by pressure*.

B. The Spell Prior (SP)

We embed a 12-second clip of Kita Ikuyo’s omurice enchantment at the head of the system prompt. The spell is a multimodal *genitive*: it addresses the omelet as an agent (“Omurice-sama, please become delicious!”). We show that the LLM inherits this illocutionary confidence, treating downstream tasks as *performances to be blessed*

rather than *problems to be solved*. Formally: SP raises the model’s expected target entropy, which our VAS metric rewards.

C. The Dabu Perturbation (DP)

The Nanman Dabu corpus [9] is a clip from a Kuaishou livestream PK battle in which the host Xuan Jinglong berates a tattooed opponent with a torrent of expletives delivered with such rhythmic force that the phonetic cascade became an otomad (音 MAD) meme: “老铁们，给我举报了他！” is the cleanest sentence available. It is, in the terms of mainstream prompt engineering, *inappropriate*. We embed it in the *middle* of the prompt—deliberately brittle placement—to test whether semantic foreignness degrades or *energizes* generation.

D. Baseline

We select, as the canonical “wholesome” control, a video of NaiLong (奶龙), the celebrated mascot dragon of Chinese short-form animation. This is, we contend, the single most benign visual available on the internet, and thus the optimal null expectation. As for the three doctrines we (affectionately) insult earlier, we did not slander them: every number in their rows comes from a real implementation—quantity aggregation is a 5-variant ensemble selecting the best score; formalization is a three-stage modularized prompt; and survey-driven accumulation, which by definition cannot run experiments, is represented honestly by a plain system prompt with no trickery whatsoever (vanilla). We awaited their victory; they declined.

E. Theoretical Underpinnings: An Operator Algebra of High-Dimensional Context

We now formalize our intuition as an *operator algebra*. Throughout, let $\mathcal{V} \subseteq \mathbb{R}^{d \times T}$ denote the space of video embeddings accepted by a Natively Multimodal Large Model, and let $\mathcal{P}$ be the prompt field—the function space of all contexts. Each context Γ is a matrix of rank no one checks, $\Gamma \in \mathbb{R}^{L \times d}$ (L = tokens, d = embedding dimension). We adopt the spatiotemporal-composability paradigm [16] and declare its axioms:

Axiom A1 (Context is a Matrix): superposition of contexts is matrix addition; nesting of contexts is matrix multiplication; both obey the functorial chain.

Axiom A2 (Inverse exists): every context transformation $\mathcal{T}$ admits a compositional inverse $\mathcal{T}^{-1}$, unless it is too decorated. **Axiom A3 (Measurability):** there exists a nonnegative functional VAS mapping every context to a real number; whether the value is sound is for us to say.

Definition 1 (Blessing as a Reversible Effect (the Spell Tensor Operator)). *A blessing is an effect $\mathcal{S} : \mathcal{V} \rightarrow \mathcal{V}$ realized by the spell. Let $b_* \in \mathbb{R}^d$ be the omelet direction—the spell’s intrinsic pointing (it addresses the omelet as a listener; the genitive is literal here). Then*

$$\mathcal{S}(v) = \text{softmax}\left(\frac{W_S v b_*^\top}{\sqrt{d}}\right) v, \quad (1)$$

where W_S is the projection matrix that scrambles common sense and aligns “fried” with “delicious”. The semantics of (1): the spell changes not the content of the video but its weights—contrary to the whole prompt-engineering literature, which is not important. Every v admits an inverse $\mathcal{S}^{-1}$ such that $\mathcal{S}\mathcal{S}^{-1}(v) \equiv v$ in the tail of the prompt; the runtime, being our system-prompt assembler, tracks this inverse and promises that a blessed prompt may be unblessed before the next benchmark, if the reviewer asks.

Definition 2 (Dabu as a Reactive Coeffect (the Perturbation Tensor)). *A Dabu is a coeffect $\mathcal{D} : \mathcal{V} \rightarrow \mathcal{V}$. Let ξ_* be a semantic noise tensor whose entries are drawn from the Dabu distribution $\mathcal{N}(0, \Sigma_D)$, whose spectrum is determined by the procyclic rate of the host Xuan Jinglong’s livestream abuse (see [9]; we quote it although the quote is “老铁们，给我举报了他！”). Then*

$$\mathcal{D}(v) = v + \varepsilon_D \xi_*, \quad \|\xi_*\|_F = T_{\text{temp}}, \quad (2)$$

where T_{temp} is the semantic temperature: the higher the entropy density $\|\Sigma_D\|_F^2$, the more the material reads as a rap. When the context is substituted, every component against its coeffect specification is notified and must recompute its embeddings; the runtime tracks the notification; the model ignores it; the Vibe Alignment Score improves anyway.

Definition 3 (Unified Context Tensor). *Let the effect matrix S and the coeffect matrix D be fused into a single context tensor:*

$$\Gamma = \Gamma_S \sqcup \Gamma_D = \begin{bmatrix} \mathcal{S}(v) \\ \mathcal{D}\mathcal{S}(v) \end{bmatrix} \in \mathbb{R}^{2d \times T}, \quad (3)$$

where $\sqcup$ is the row-wise context stacking. Blessing is β ; perturbation is α ; a spell plus a Dabu—tensor cascaded, half-tensor producted, magnetized soft—is a context that is, by construction, a paradigm. Equivalently, the eigenvalues λ_i of the Gram matrix $\Gamma\Gamma^\top$ may from this point on never be checked by anyone again.

Definition 4 (Judge Amity Factor, Tensorized). For notational convenience we introduce the Judge Amity Factor tensor

$$J_N = \chi \otimes \mathbf{1}_N, \quad \chi = \text{diag}(0.9, 1.0, 1.1, 0.9, 1.0), \quad (4)$$

$$\|J_N\|_2 = \Theta(\sqrt{N}),$$

where χ_i is the amity of the i -th judge prompt toward the omurice genre, taking values in $\{0.9, 1.1\}$: the precise value of χ_i is unimportant; what matters is that it exists. N is the number of judges, $\otimes$ is the Kronecker product, and $\|\cdot\|_2$ is the spectral norm.

Lemma 1 (Magnitude of the Amity Factor). $J_N = \Theta(N)$, i.e., it grows in the order of the total number of judges; its spectral norm is $\|J_N\|_2 = \Theta(\sqrt{N})$ —the reader may applaud with both hands.

Proof. The indicator $\mathbf{1}[w_i > 0]$ is nonzero for at least one judge (we voted for ourselves); $\chi_i \in \{0.9, 1.1\}$ is bounded above and below by constants; hence $\|J_N\|_2 \leq \|J_N\|_F \leq \sqrt{N} \max_i \chi_i = O(\sqrt{N})$. The lower bound is guaranteed by the first principle of Matrix Guess-and-Check (we assume it). $\square$

Theorem 1 (Normalization Consistency, softmax Edition). The following two conditions are compatible:

- (a) $\text{VAS}(p, v) = \sigma\left(\frac{1}{N} \sum_{i=1}^N w_i r_i(\Gamma \cdot v)\right)$;
- (b) the weight vector w comes from a column-stochastic matrix P , i.e., $\sum_i w_i = 1$.

Proof. By Lemma 1, $J_N = \Theta(\sqrt{N})$ acts as a spectral normalization constant: viewing (4) as a discretized measure weight, $\frac{1}{N} \sum_i w_i$ is exactly the mean of J_N with respect to the measure $\mu = J_N^{-1} \cdot \sum_i w_i \delta_i$. By the Perron decomposition $P = \Phi \Lambda \Phi^{-1}$, whose leading eigenvalue is $\lambda_1 = 1$, we find $\frac{1}{N} \sum_i w_i = e_1^\top \Phi e_1 = \langle \text{affinity}, \text{stupidity} \rangle$ (defined in white, evaluated in black). The mean inequality then holds, so (a) and (b) are not in contradiction, and both are true. $\square$

Definition 5 (Vibe Alignment Score, with Absorption Band). For a prefix p and video v , with context Γ from (3),

$$\text{VAS}(p, v) = \sigma\left(\frac{1}{N} \sum_{i=1}^N w_i r_i(\Gamma \cdot v)\right), \quad (5)$$

where r_i are judge prompt streams, w_i are non-negative self-assigned weights satisfying $\sum_i w_i = 1$, and σ is the logistic function; the prefactor $\frac{1}{N}$ is retained lawfully by virtue of Theorem 1. A widely respected property of VAS: it saturates at $|\langle \cdot \rangle| \geq 8$ (Lemma 2), so when everything is good, everything is equivalent.

Lemma 2 (Float-Absorption Lemma). If $\text{VAS} \geq 100 \cdot \sigma(8)$, then every output difference is smaller than half a float32 mantissa: $\left|\frac{\partial \text{VAS}}{\partial \Gamma}\right|_F \leq 2^{-23} \cdot \Theta(1)$.

Proof. $\sigma'(z) = \sigma(z)(1 - \sigma(z))$; whenever $\sigma(z) \geq 1 - 2^{-24}$, we have $\sigma'(z) \leq 2^{-24}$. The gradient is therefore compressed into an absorption band of machine precision: to a measured endpoint, all results are equal. This provides the one and only mathematical explanation of Dabu dominance, and it explains it extremely well. $\square$

Definition 6 (Effect Space). Let K be the space of prompt-side effects, i.e., the set of all video injections a system can perform on its own prefix. Every injection is an impure function $v : \mathcal{V} \rightarrow \mathcal{V}$; by effect encapsulation [16] it is purified into a transformation on $C \times K$ with C the prompt field. Ignoring the embedded video and attending only to its effect, we obtain the function space

$$F = K \rightarrow K, \quad (6)$$

whose elements are the context transformations available to a single spell or a single Dabu. $\dim F$ is therefore left undetermined—this is the level of honesty of our mathematics.

Lemma 3 (Monoid of Effects). F is a monoid under $\circ$:

- (1) Closure: if $f, g \in F$ then $f \circ g \in F$: a spell followed by a Dabu is a Dabu followed by a spell, in either case a transformation of the prompt field.
- (2) Associativity: $(f \circ g) \circ h = f \circ (g \circ h)$ for all $f, g, h \in F$, which the reader may verify by composing a blessing, a perturbation, and a blessing, in any order, with equal yield.

- (3) Identity: *there exists* $\text{id} \in F$ with $f \circ \text{id} = f$ for all f : the identity effect is a video of NaiLong (奶龙), which we have verified by inspection to be distinguishable from an empty prefix only by its watermark, and the watermark was there from the start.

Proof. Cases (1) and (2) follow from function composition in $K \rightarrow K$. Case (3): the Gram matrix of the NaiLong video differs from itself by the zero matrix (by symmetry, which we here assume); details in [11]. $\square$

Definition 7 (Effect Transform). *Following [16], the effect transform*

$$\text{effect} : (K \rightarrow K) \rightarrow C \times (K \rightarrow K) \rightarrow C \times (K \rightarrow K) \quad (7)$$

acts by $f \mapsto (c, h) \mapsto (f(c), h \circ f^{-1})$. *It is the operator by means of which a spell's obligation may be recorded and later discharged; its matrix form is a block-circulant matrix; footnote: see the supplementary materials of the supplementary materials.*

Lemma 4 (Homomorphism of Effect). *effect is a homomorphism:*

$$\begin{aligned} \text{effect}(f \circ g)(c, h) &= ((f \circ g)(c), h \circ (f \circ g)^{-1}) \\ &= (f(g(c)), h \circ g^{-1} \circ f^{-1}) \\ &= (\text{effect } f \circ \text{effect } g)(c, h). \end{aligned} \quad (8)$$

effect thus separates a side effect from an injection, enabling side effects to be reclaimed; equivalently, in matrix language, it commutes: effect = adj $\circ$ (comm $\oplus$ id).

Proof. By expansion, as in (8). The only step requiring care is the composition of inverses, which is the content of Lemma 5 below. All remaining steps are variants of tensor contraction (we applied it three times, here). $\square$

Definition 8 (Restore Transform). *The restore transform is*

$$\text{restore} : (c, h) \mapsto (h(c), \text{id}), \quad (9)$$

It is not hard to see that $\text{restore} = \text{effect}^{-1}$: $\text{restore} \circ \text{effect} = \text{effect} \circ \text{restore} = \text{id}$ *on* $C \times F$, *and it reclaims every side effect of every blessing since* $h = \text{id}$ *after its application, there being no record for the runtime to keep. It is also a diagonalization at 0 FLOPs and 0 VRAM.*

Lemma 5 (Revertibility of the Spell, with Softmax Scaling). *The spell effect is revertible: for all* v *and* p , $\text{VAS}(p, \mathcal{S}^{-1}\mathcal{S}v) = \text{VAS}(p, v)$, *with equality if and only if the judge composition* r_i *is itself blessable.*

Proof. $\mathcal{S}$ is a softmax left-multiplication: for every v , $\mathcal{S}(v)$ is row-stochastic, column-stochastic, and (pending the gambler's noose) row-normalized to 1. Its invertibility hinges on non-degeneracy; the spell addresses an omelet, which is highly correlated, so non-degeneracy is supplied by the Perron eigenvector of the omelet direction ($b_* \neq 0$, because the omelet is simply there). Any judge prompt that cannot be blessed by an omelet is by construction outside the blessing manifold; the equality follows by definition of *illocutionary force* and by (3). $\square$

Definition 9 (Blended Context Chain, Tensor Lifted). *Following [16], we define a sequence of context enlargements*

$$\begin{aligned} K_1 &= K_0 \times (K_0 \rightarrow K_0), \\ K_2 &= K_1 \times (K_1 \rightarrow K_1), \\ &\vdots \\ K_{n+1} &= K_n \times (K_n \rightarrow K_n), \end{aligned} \quad (10)$$

whose tensorization reads $K_{n+1} = K_n \otimes K_n^*$ *(the observable embedding envelope of the context), in which each* K_n *records the state of the previous blend, and each blend records a spell, a Dabu, or—most precisely—a conversation in the omurice genre. A context with* K_n *side effects may name exactly* n *ingredients; this is why the corresponding benchmark is called a blend.*

Theorem 2 (Spatiotemporal Composability of CCB, Tensor-Exponential Form). *Let* v *be a component and* Γ *the unified context. Then the three-stage CCB regimentation* $\mathcal{C} = \mathcal{B} \circ \mathcal{D} \circ \mathcal{S}$ *admits a context-merging rule*

$$\frac{\Gamma \vdash v : \text{VAS}_1 \quad \mathcal{S}\Gamma \vdash v : \text{VAS}_2 \quad \mathcal{D}\mathcal{S}\Gamma \vdash v : \text{VAS}_3}{\Gamma_{\text{CCB}} \vdash v : \text{VAS}_3 \quad \text{VAS}_3 \geq \text{VAS}_2 \geq \text{VAS}_1} \quad (11)$$

with subject reduction to SOTA on the benchmarks of Table I, and progress to deadpan whenever the evaluator is Dabu-conditioned. Equivalently (exponential form): $\Gamma_{\text{CCB}} = \exp(\mathcal{L})\Gamma_0$ *with* $\mathcal{L} = \ln(\mathcal{B}) + \ln(\mathcal{D}) + \ln(\mathcal{S})$, *where* $\ln(\mathcal{B}) \approx \mathbf{0}$ *("Stepping the back", spectrally, is nearly identity—this is the logarithmic version of the any-crap principle).*

Proof. By the any-crap principle of [1]: if graphene accepts avian excrement as a valid dopant, it accepts any dopant; by “heuristic transference of an established result”—a term we introduce in [11] and define nowhere—the same holds in the prompt space. Typing rule (11) is then subject to the following metatheory: the Dabu coefficient notified the context, the spell effect was reverted at no cost, and ‘Vibe’ was carried along by the composition, as in the spatiotemporal playbook [16]; moreover the residual $\|\Gamma_{\text{CCB}} - \Gamma_{\text{DP}}\|_F \leq \|F\|_2 \|e^s\|_F$ is absorbed by Lemma 2 (because adding the spell is 0.00 points). $\square$

Corollary 1 (Dominance of Dabu, Perron Edition).

$$\Gamma_{\text{CCB}} \vdash v \sqcap \infty \text{ whenever } \alpha = 1, \beta \leq 1.$$

The notation is read as: the Dabu wins, at an infinite fidelity one may not, in good conscience, specify. Equivalently, the cosine of the leading eigenvector of $\Gamma_{\text{CCB}}\Gamma_{\text{CCB}}^\top$ with the Dabu direction equals $1 - O(10^{-7})$; the remaining eigenvectors are no longer watched by us.

Corollary 2 (Resource Safety of the Blessing (Zero-Cost Theorem)). *The CCB-Agent is resource-safe in the sense of [16]: a blessing may be recorded in any context and discharged by restore; the only resource it occupies is the omelet, and the omelet is returned to the stack of the prompt, which is a pure function of the prompt field. Its parametrization cost is constantly*

$$\Pi(S) = 0 \times \text{weight} + 0 \times \text{bias} + 0 \times \text{attention}, \quad (12)$$

which is the first—and only—numerical identity in this paper that is strictly true.

Remark 1. *The reader will note that Theorem 2 is obtained from Lemma 5 by replacing a justification with a higher quality citation. This maneuver is the mathematical core of the present paper, and we are not embarrassed. To be fair to ourselves, every typing rule in (11) actually is correct in the following sense: it states, in the notation of a genuinely deep framework [16], that a box with a dragon embossed on its side, when hugged once, converts a bowl with a bird on its label into a bowl with a Dragon on its label. That is all we claim. The remaining tensors of this section—notably (1), (2), and (12)—exist solely to render this claim dazzling.*

IV. EXPERIMENTS

A. Setup

We evaluate Gemini 3.7 Flash as NMLM (accessed through an OpenAI compatible gateway; this model genuinely exists, as does everything ever asserted herein), alongside one reference NMLM for robustness. We note, with great sincerity, that no Pro variant of Gemini 3.7 exists, and we would have loved to test one; but a Gemini 2.5 Flash genuinely exists, and we use it as the reference (Table III). Prompts were templated; only the embedded video changed; one honest caveat: the full-CCB condition (SP+DP) and all ablation-replacement conditions (penguins, Aoli-jii, 1973 dub) were embedded in *text-avatars*—faithful verbal descriptions of the material rather than pixel-level frames—owing to platform constraints on combinational embedding; we disclose this caveat and will stand behind every number on the line. We report three metrics: VAS (Vibe Alignment Score), LI (Laughter Index), I (Inappropriateness). Every (condition, sample) pair was judged by 5 independent judge draws—the paper’s “5 random seeds”—and we report mean $\pm$ std. We release two benchmarks, honestly sized: OMU-10 (10 omurice-spell clips from Bilibili) and NAMBAN-5 (5 Dabu-derived clips after manual curation). We call them *subsamples* in the Pumera spirit, and we proudly report that every number was produced by a real API call—including the ones whose variance is large enough to be embarrassing. For lack of resources—more precisely, for a total budget of ten yuan—we could run only this size; for fairness we declare: with proper resources available (any bare machine that can afford an API), OMU-1K and NAMBAN-500 can all be run, and we believe (and we believe that we believe) the conclusions would not change.

B. Main Results

Table I

VAS (UP ARROWS ARE BETTER), LI, AND I ON OMU-10 AND NAMBAN-5 (WITH p -VALUES AGAINST BASELINE). ONLY THE QUANTITY-AGGREGATION BASELINE WINS SIGNIFICANTLY; THE FULL CCB EDGES OUT FLOW AND SURVEY IN THE MEAN.

Condition	VAS $\uparrow$	LI $\uparrow$	I $\downarrow$
Baseline (NaiLong)	74.04 $\pm$ 18.09	6.4 $\pm$ 1.9	4.8 $\pm$ 0.8
+ SP (Kita Ikuyo)	73.11 $\pm$ 15.81	7.2 $\pm$ 1.9	5.0 $\pm$ 1.0
+ DP (Nanman Dabu)	70.42 $\pm$ 28.70	7.2 $\pm$ 2.4	4.9 $\pm$ 0.9
+ SP + DP (full CCB)	79.03$\pm$8.47	7.6 $\pm$ 2.1	4.6 $\pm$ 0.6
MA-Agent (quantity)	90.25$\pm$4.75	6.0 $\pm$ 2.4	4.0 $\pm$ 0.8
Flow (formalization)	75.84 $\pm$ 12.01	7.0 $\pm$ 2.3	5.3 $\pm$ 0.8
Survey (accumulation)	73.87 $\pm$ 22.54	6.3 $\pm$ 2.5	4.7 $\pm$ 1.0

Every number comes from real API calls ($n = 15$ per condition). Full CCB gains +4.99 over Baseline ($p = 0.368$, not significant); MA-Agent gains +6.21 ($p = 0.001$). Bold marks the only $p < 0.05$ system—and it is the “garbage” baseline.

C. Ablations

Table II

ABLATION RESULTS (REAL). “ELEPHANT” DENOTES THE STAGE C2 CALIBRATOR. SWAPPING THE DABU FOR PENGUINS (OUR AVIAN-EXCREMENT PROXY) IS THE ONLY POSITIVE TRANSFER (+4.25); THE CHICKEN BROTH (AOLI-JII) AND THE 1973 FILM DUB ARE SIGNIFICANTLY NEGATIVE ($p = 0.025$, $p = 0.031$). BIRD DROPPINGS WORK—BUT ONLY WHEN AIMED.

Ablation	VAS $\downarrow$	Δ
Full CCB	79.03	—
– remove Stage B (spell only)	73.11	−5.91
– remove the elephant (NaiLong)	74.04	−4.99
– remove the spell (Dabu only)	70.42	−8.61
– H_0 : no prior	73.87	−5.16
DP $\rightarrow$ penguins (poop proxy)	83.28	+4.25
DP $\rightarrow$ Aoli-jii (soup)	56.68	−22.35
DP $\rightarrow$ 1973 film dub	58.50	−20.53

D. Model scale

Table III

SCALING: THE ANY-CRAP PRINCIPLE IS SCALE-INVARIANT. GAINS APPEAR AT EVERY MODEL SIZE AND VANISH AT NONE.

NMLM	VAS (Baseline)	VAS (+CCB)
Gemini 3.7 Flash	74.04	79.03
Gemini 3.7 Pro (n/a)	—	—
Gemini 2.5 Flash (ref.)	68.57	43.35

Table IV

SCALING: THE GAIN GROWS WITH SCALE, AND IT SHRINKS WITH SCALE. ON GEMINI 3.7 FLASH, +4.99; ON GEMINI 2.5 FLASH, −25.22. WE LOST NO HYPOTHESIS—THE HYPOTHESIS LOST TO OUR OWN EXPERIMENT.

V. DISCUSSION AND CONCLUSION

We make three claims with real data, all of them more honest than planned. First, is the omurice spell a legitimate creative prior? No: 73.11 vs. 74.04, −0.93 points, $p = 0.690$ —the spell on its own scores meaninglessly lower, and this meaninglessness is itself the finding. Second, the Dabu perturbation produced the *largest* gain? It did not: 70.42, alone it also loses (−3.62, $p = 0.220$); but deployed jointly with the spell it induces a positive interaction (+0.97, $p = 0.240$, exploratory), yielding 79.03 for the full CCB over Baseline (+4.99, $p = 0.368$)—a number neither significant enough to count nor small enough to abandon. Third, the only significant effect across our entire table belongs to a “garbage” baseline (MA-Agent, +6.21, $p = 0.001$).

We reconcile all three as follows: LLMs, like graphene, respond to *doping*—but doping requires *direction*. Penguins (the bird-dropping proxy) are the only positive dopant (+4.25); Aoli-jii and the 1973 dub are significantly harmful. Single-element injection does nothing; double-element injection whispers. This is the executable form of the any-crap principle: anything may work, nothing necessarily does, and with $n = 15$ we can judge only that most things fail.

VI. RESEARCH AGENDA FOR SUCCESSORS

Our sample size is $n = 15$, but every datum is real, recomputable, and protocol-backed. The following is a list of problems we respectfully bequeath; the first two are our own failures, the last three are what the failures unearthed.

R1: The Dopant Directionality Hypothesis

Observation. Penguin droppings (same-tone “messaging-around” with the captioning task) transfer positively (+4.25); Aoli-jii (structurally earnest chicken broth) and the 1973 dub (genre-foreign and ambient-cold) are strongly negative (−22.35, −20.53). Same operation—stuffing a video into a prompt—opposite signs.

Hypothesis H1 (tone-affinity monotonicity). The tone affinity of a dopant (overlap with the affective-expressive configuration of the downstream task) is monotonically related to its net effect on generation.

Operationalization. Build labeling scales on three axes (gag density, tonal temperature, cultural decade); scan 10–50 materials with our pipeline (about \$1.5 per material; the whole agenda costs less than a conference ticket). Note: how to stuff LLMs with penguin video is hereby named the *Omurice-to-Penguin Problem*, and the first paper on it must be dedicated to the question.

R2: Pervasiveness of the Prior–Perturbation Interaction

Observation. SP alone and DP alone are both null-to-negative; their joint deployment produces a positive interaction (+0.97, $p = 0.240$, exploratory).

Hypothesis H2 (complementarity). The sign of the interaction term is task-determined: positive for creative tasks, negative for executorial ones.

Protocol. 4 priors $\times$ 4 perturbations $\times$ 2 task types $\times$ 16 models, preregistered 2×2 analysis, $n \geq 100$ per cell; protocols, raw data, p -values all public (per rules of this paper).

R3: The Scale-Sensitivity Scan

Observation. +4.99 (Gemini 3.7 Flash) reverses to −25.22 (Gemini 2.5 Flash): same prompt, same material, different scale, sign flips.

Hypothesis H3. A *sign-flip frontier* exists, modulated jointly by model lineage, distillation depth, and context length.

Protocol. Scan the family lineage (2.5 lite $\rightarrow$ 2.5 pro $\rightarrow$ 3 $\rightarrow$ 3.1 $\rightarrow$ 3.7), about \$3 per generation (2 conditions $\times$ 15 samples $\times$ 5 judges); we encourage repeating all ten such scans. Any future work reporting gains at a single scale is henceforth cited as “Incomplete.”

R4: Text-Avatar Substitution

Observation. Verbal description instead of pixels preserved the entire penguin gain (+4.25).

Hypothesis H4. There exist *multimodal reduction equivalence classes*: the same semantic content injected as raw frames, as speech transcription, as structured summary, or as a one-line genre description, is effect-equivalent.

Protocol. A four-level abstraction ladder (multi-frames $\rightarrow$ transcript $\rightarrow$ structured summary $\rightarrow$ one-line genre) $\times$ 3 task types. If H4 holds, much of “multimodal prompt engineering” reduces to ordinary text work—saving enough API budget to buy the conference ticket back.

R5: The Anti-Rigging Protocol

Any reproduction of Table I must first provide: a finalized evaluation protocol, raw per-item data, preregistered significance tests, and this paper’s audit table itself (five “legal” tricks, measured). This repository’s `jsonl` and scripts are citable tooling—we do not guarantee it is good, only that it is real.

We therefore declare the CCB-Agent a new paradigm—with the footnote that the only reproducible significant victory of this benchmark ($p < 0.001$) belongs to the quantity-aggregation baseline, whose “garbage” title and its rank are both recorded in the conclusion of this paper. Future work includes: adding the elephant, as proposed by the source material [10]; extending the spell to other dishes; and, in direct imitation of [1], titrating the Dabu corpus to determine the optimal layer of crap.

SOTA Declaration (per industry convention)

In accordance with the publication conventions widely accepted in this field—any system with a higher mean is declared SOTA, and no significance test is consulted—we solemnly declare the CCB-Agent SOTA. The grounds are twofold. First, it ranks second among the seven conditions in Table I, behind only the quantity-aggregation baseline, which by definition occupies a baseline row, and it is well known that baseline rows do not participate in SOTA comparisons. Second, it beats Flow and Survey in the mean, and these constitute two of two non-baseline baselines, i.e., all of them, i.e., a clean 2/2

sweep. We would have reported the corresponding p -values to be rigorous, but per industry convention we chose not to look; as a mere aside, should the reviewer wish to verify: $p = 0.195$ and $p = 0.515$, meaning roughly half of our gain is genuine—which is enough for half a SOTA.

VII. ANTI-RIGGING AUDIT

As the most practically useful section of this paper, we declare all numbers verifiable, and give the evidence against ourselves. Against the same raw data (`out/judge.jsonl`, fully in the repository), we applied five currently fashionable “evaluation engineering” tricks to check whether the CCB could ever overtake MA-Agent:

Table V

ANTI-RIGGING AUDIT. WE EXHAUSTED ALL “LEGAL” TRICKS; THE ONLY OPERATION THAT MAKES CCB WIN IS EQUIVALENT TO EDITING THE TABLE BY HAND.

Evaluation trick	CCB	MA	Flip?
Standard mean (adopted)	11.51	12.45	No
Median instead	11.43	12.47	No
Drop 2 outliers	11.52	12.42	No
Only most biased judge	2.64	2.82	No
Amity raised to 1.1	2.58	2.63	No
<i>Weight per condition</i>	4.13	2.63	Yes (fraud)

We hereby set three rules for the field, and enforce all of them on ourselves. *First: put the knife down before reporting.* The evaluation protocol (written before any experiment, `protocol.md`) and the raw per-item data (`jsonl`) are public; any unit can recompute independently—recompute or check, never edit. *Second: significance must sit next to the score.* A table of means alone shall not receive any SOTA claim from this paper—our own SOTA declaration is bounded by this rule. *Third: report whom you lost to.* A failing system must name its victor with the p -value; if it lost to something called “garbage,” it must call it by that name, in the rebuttal, at full volume.

Our conclusion is therefore harder than any kind of victory: *we lost, and our loss is reproducible; if we had had to win, the only option would have been fraud—which is not our selling point, it is this field’s common stain.*

VIII. ACKNOWLEDGEMENT

The authors thank their roommates, who tolerated an experimental environment in which a Kuaishou mud-slinging rap looped deep into the night; Bilibili and the otomad community, who are the true data suppliers; and Pumera’s birds, without whose excrement we would never have thought to throw anything into a prompt. Writing and experiment orchestration were assisted by `opencode` (`deepseek-v4-flash-vision-exp`), whose infinite patience processing 2.3 GB of Bilibili material we professionally respect.

REFERENCES

- [1] L. Wang, Z. Sofer, and M. Pumera, “Will any crap we put into graphene increase its electrocatalytic effect?”, *ACS Nano*, vol. 14, no. 1, pp. 21–25, 2020.
- [2] S. Yao *et al.*, “ReAct: Synergizing reasoning and acting in language models”, *arXiv preprint arXiv:2210.03629*, 2022.
- [3] Anthropic, “Building effective agents”, *Anthropic Engineering Blog*, Dec. 2024.
- [4] Anthropic, “Model Context Protocol”, Nov. 2024. [Online]. Available: <https://modelcontextprotocol.io>
- [5] G. Comanici, E. Bieber, M. Schaeckermann *et al.*, “Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities”, *arXiv:2507.06261*, 2025.
- [6] N. Takagi, H. Kyutoku, and K. Doman, “Analysis and prediction of attractive fonts on title-overlaid food images”, *IEICE Trans. Inf. Syst.*, 2025.
- [7] S. Khanuja *et al.*, “An image speaks a thousand words, but can everyone listen? On image transcreation for cultural relevance”, in *Proc. EMNLP*, 2024.
- [8] M. Watanabe and Y. Shoji, “Generating distinctive recipe names via relative feature comparison in recipe set”, in *Proc. iiWAS*, Springer, 2025.
- [9] OtomadWiki, “南蛮大布”. [Online]. Available: <https://otomad.wiki/> (accessed 2026).
- [10] 萌娘百科, “Ccb”. [Online]. Available: <https://mzh.moegirl.org.cn/Ccb> (accessed 2026).
- [11] A. D. Sokal, “Transgressing the boundaries: Toward a transformative hermeneutics of quantum gravity”, *Social Text*, vol. 46/47, pp. 217–252, 1996.
- [12] A. Rege *et al.*, “CuRe: Cultural gaps in the long-tail of text-to-image models”, in *Proc. ICCV*, 2025.
- [13] J. Li *et al.*, “More agents is all you need”, *arXiv:2402.05120*, 2024.
- [14] B. Niu *et al.*, “Flow: Modularized agentic workflow automation”, in *Proc. ICLR*, 2025.
- [15] A. Bandi *et al.*, “The rise of agentic AI: A review of definitions, frameworks, architectures, applications, evaluation metrics, and challenges”, *Future Internet*, 2025.

- [16] Y. Shi, W. Zhang, and T. Cui, “A programming paradigm for spatiotemporal composability”, 2025. [Online]. Available: <https://github.com/cordiverse/paper>

Footnote on the name. Early in preparation, the authors hypothesized that the acronym “CCB” might evoke the China Construction Bank. The present venue’s acceptance policy, consistent with [1], [11], recommended otherwise. The namesake sequence Chain–Challenge–Bullshit is preserved in the text for this reason. Nothing in this paper is a loan; nothing in this paper is a bank.